

ANN

December 4, 2024

```
[86]: # import necessary libraries
import pandas as pd
import numpy as np
```

```
[74]: data = pd.read_csv('df_merged_weekly.csv', index_col=0)
data['dt'] = pd.to_datetime(data['dt'])
data.head()
```

```
[74]:
```

| | week | dt | Cases | temp | feels_like | pressure | humidity | \ |
|---|--------|------------|-------|--------|------------|----------|----------|---|
| 0 | 2022-1 | 2022-01-01 | 7.0 | 298.38 | 299.64 | 1013.0 | 87.0 | |
| 1 | 2022-2 | 2022-01-08 | 1.0 | 302.23 | 308.88 | 1011.0 | 68.0 | |
| 2 | 2022-3 | 2022-01-15 | 2.0 | 300.74 | 304.09 | 1013.0 | 58.0 | |
| 3 | 2022-4 | 2022-01-22 | 3.0 | 302.42 | 309.89 | 1010.0 | 63.0 | |
| 4 | 2022-5 | 2022-01-29 | 3.0 | 303.37 | 308.69 | 1010.0 | 64.0 | |

| | precipitation |
|---|---------------|
| 0 | 3.14 |
| 1 | 0.00 |
| 2 | 0.00 |
| 3 | 0.00 |
| 4 | 0.42 |

```
[75]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 140 entries, 0 to 139
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   week            140 non-null    object
1   dt              140 non-null    datetime64[ns]
2   Cases           140 non-null    float64
3   temp            140 non-null    float64
4   feels_like      140 non-null    float64
5   pressure        140 non-null    float64
6   humidity        140 non-null    float64
7   precipitation    140 non-null    float64
```

dtypes: datetime64[ns](1), float64(6), object(1)
memory usage: 9.8+ KB

1 Columns

| column | description | type |
|---------------|---|----------|
| week | week corresponding to the specific year | datetime |
| dt | specific date | datetime |
| Cases | Dengue Cases | integer |
| temp | Average temperature in the given week | numeric |
| feels_like | Average feels_like in the given week | numeric |
| pressure | Average pressure in the given week | numeric |
| humidity | Average humidity in the given week | numeric |
| precipitation | Total precipitation in the given week | numeric |

1.1 Training data

Data used to be trained range from 2022 to 2023

1.2 Testing data

Data used to be tested is 2024 (also used for training)

```
[76]: # Feature Engineering
from sklearn.preprocessing import MinMaxScaler
standard_scaler = MinMaxScaler()
features = ["temp", "humidity", "precipitation"]
X = data[features]
X = standard_scaler.fit_transform(X)
temp = X[:,0]
humidity = X[:,1]
precipitation = X[:,2]
data['temp'] = temp
data['precipitation'] = precipitation
data['humidity'] = humidity
```

```
[77]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset
from sklearn.metrics import mean_absolute_error, mean_squared_error
class Model(nn.Module):
    def __init__(self, input_size, hidden_sizes, output_size, dropouts):
        super().__init__()
        self.layer1 = nn.Linear(input_size, hidden_sizes[0])
        self.layer2 = nn.Linear(hidden_sizes[0], hidden_sizes[1])
        self.layer3 = nn.Linear(hidden_sizes[1], hidden_sizes[2])
```

```

        self.output = nn.Linear(hidden_sizes[2], output_size)

        self.dropouts = nn.ModuleList([nn.Dropout(dropout) for dropout in
→dropouts])
        self.relu = nn.ReLU()

    def forward(self, x):
        x = self.relu(self.layer1(x))
        x = self.dropouts[0](x)
        x = self.relu(self.layer2(x))
        x = self.dropouts[1](x)
        x = self.relu(self.layer3(x))
        x = self.dropouts[2](x)
        x = self.output(x)
        return x

def train_ANN(data, features, target, date_col, n_ahead):
    # Select relevant columns
    rel_col = features + [target] + [date_col]
    data = data[rel_col]

    # Getting the first and last index for the year 2024
    first_2024 = data[data[date_col].dt.year == 2024].index[0]
    last_2024 = data[data[date_col].dt.year == 2024].index[-1]

    first_2024subn = first_2024 - n_ahead
    last_2024subn = last_2024 - n_ahead

    # Define lags
    env_lags = [1,2,3,4] # 2-week lag for environment features
    cases_lag = range(1, 12) # 1 to 11 week lags for target variable

    # Create lagged features for environment and target variables
    for lag in env_lags:
        for feature in features:
            data[f'{feature}_lag_{lag}'] = data[feature].shift(lag)

    for lag in cases_lag:
        data[f'{target}_lag_{lag}'] = data[target].shift(lag)

    # Remove any rows with missing values due to lagging
    data = data.dropna()

    # Initialize list for storing predictions
    predict_data = []

    # Loop through data to get n-week ahead prediction

```

```

index_begin = first_2024subn
index_end = last_2024subn

while index_begin <= index_end:
    # Split data into training and testing sets
    train_data = data[data.index <= index_begin]
    test_data = data.iloc[data.index.get_loc(index_begin):data.index.
↪get_loc(index_begin) + n_ahead]

    X_train = train_data.drop(columns=[target, date_col])
    y_train = train_data[target]
    X_test = test_data.drop(columns=[target, date_col])
    y_test = test_data[target]

    # Convert to PyTorch tensors
    X_train = torch.tensor(X_train.values, dtype=torch.float32)
    y_train = torch.tensor(y_train.values, dtype=torch.float32)
    X_test = torch.tensor(X_test.values, dtype=torch.float32)
    y_test = torch.tensor(y_test.values, dtype=torch.float32)

    # Combine inputs and labels into a Dataset
    train_dataset = TensorDataset(X_train, y_train)
    train_loader = DataLoader(train_dataset, batch_size=32, shuffle=False)

    # Initialize the ANN Model
    input_size = X_train.shape[1]
    hidden_sizes = [48, 32, 19]
    dropouts = [0.3, 0.2, 0.1]
    output_size = 1
    model = Model(input_size, hidden_sizes, output_size, dropouts)

    # Define optimizer and loss function
    optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
    loss_fn = nn.MSELoss()

    # Train the model
    model.train()
    num_epochs = 1000
    for epoch in range(num_epochs):
        for batch_X, batch_y in train_loader:
            optimizer.zero_grad()

            # Get model predictions (assuming output size is n_ahead)
            predictions = model(batch_X)

            # Compute loss (MSE)
            loss = loss_fn(predictions.squeeze(), batch_y)

```

```

        # Backpropagation
        loss.backward()
        optimizer.step()
        #print(f'Epoch {epoch+1}/{num_epochs} has passed')

    # Evaluate on test data
    model.eval()
    with torch.no_grad():
        predictions = model(X_test).squeeze().numpy()
    try:
        predictions = predictions[-1]
    except IndexError:
        pass
    predict_data.append(predictions)
    index_begin += 1

# Calculate Mean Absolute Error
actual_data = data[data[date_col].dt.year == 2024][target].values
MAE = mean_absolute_error(actual_data, predict_data)
MSE = mean_squared_error(actual_data, predict_data)
return predict_data, MAE, MSE

```

```

[78]: # Perform Initial Testing
target="Cases"
features=["temp","humidity","precipitation"]
date = 'dt'
prediction_1_week, MAE_1_week, MSE_1_week = train_ANN(data, features, target,
↪date, 1)
prediction_4_week, MAE_4_week, MSE_4_week= train_ANN(data, features, target,
↪date, 4)

```

/tmp/ipykernel_5158/3166879319.py:46: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
data[f'{feature}_lag_{lag}'] = data[feature].shift(lag)
```

/tmp/ipykernel_5158/3166879319.py:46: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
data[f'{feature}_lag_{lag}'] = data[feature].shift(lag)
```

/home/miniloda/Documents/GitHub/Math-198.1---Special-

```
Problem/venv/lib/python3.10/site-packages/torch/nn/modules/loss.py:608:
UserWarning: Using a target size (torch.Size([1])) that is different to the
input size (torch.Size([])). This will likely lead to incorrect results due to
broadcasting. Please ensure they have the same size.
```

```
    return F.mse_loss(input, target, reduction=self.reduction)
/tmp/ipykernel_5158/3166879319.py:46: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
    data[f'{feature}_lag_{lag}'] = data[feature].shift(lag)
/tmp/ipykernel_5158/3166879319.py:46: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
    data[f'{feature}_lag_{lag}'] = data[feature].shift(lag)
/home/miniloda/Documents/GitHub/Math-198.1---Special-
Problem/venv/lib/python3.10/site-packages/torch/nn/modules/loss.py:608:
UserWarning: Using a target size (torch.Size([1])) that is different to the
input size (torch.Size([])). This will likely lead to incorrect results due to
broadcasting. Please ensure they have the same size.
    return F.mse_loss(input, target, reduction=self.reduction)
```

```
[87]: print(f'The Mean Absolute Error of the 1-Week Ahead Model is {MAE_1_week}')
      print(f'The Mean Absolute Error of the 4-Week Ahead Model is {MAE_4_week}')
      print(f'The Root Mean Squared Error of the 1-Week Ahead Model is {np.
        ↳sqrt(MSE_1_week)}')
      print(f'The Root Mean Squared Error of the 4-Week Ahead Model is {np.
        ↳sqrt(MSE_4_week)}')
```

```
The Mean Absolute Error of the 1-Week Ahead Model is 12.189200864897835
The Mean Absolute Error of the 4-Week Ahead Model is 17.1717133919398
The Root Mean Squared Error of the 1-Week Ahead Model is 23.6536483968945
The Root Mean Squared Error of the 4-Week Ahead Model is 34.12524544437012
```

```
[80]: prediction_1_week[:5]
```

```
[80]: [array(14.844003, dtype=float32),
      array(13.386841, dtype=float32),
      array(12.713357, dtype=float32),
      array(11.87863, dtype=float32),
      array(15.817395, dtype=float32)]
```


```
[81]: prediction_4_week[:5]
```

```
[81]: [np.float32(25.019201),
      np.float32(25.40394),
      np.float32(11.99758),
      np.float32(12.008745),
      np.float32(11.9914)]
```

2 Visualization


```
[82]: import matplotlib.pyplot as plt
first_2024 = int((data['dt'].dt.year==2024).idxmax())
last_2024 = int(data.loc[data['dt'].dt.year == 2024].index[-1])
date = data.loc[first_2024:last_2024]['dt']
```

```
[83]: fig, ax = plt.subplots(1,2, figsize=(15,5))
ax[0].plot(date,prediction_1_week)
ax[0].set_title('Prediction For 1-week ahead Model')
ax[1].plot(date,prediction_4_week)
ax[1].set_title('Prediction For 4-week ahead Model')
plt.suptitle('Predictions For n-week ahead Models')
plt.show()
```




```
[84]: actual_2024 = data[data['dt'].dt.year == 2024]
plt.plot(actual_2024['dt'], actual_2024['Cases'])
plt.plot()
plt.title('Actual Cases')
```

```
[84]: Text(0.5, 1.0, 'Actual Cases')
```



```
[88]: fig, ax = plt.subplots(1,2, figsize=(20,5))
ax[0].plot(data['dt'], data['Cases'], label = 'Actual Cases')
ax[0].plot(date, prediction_1_week, color = 'red', label = '1-week ahead')
ax[0].plot(date, prediction_4_week, color = 'green', label = '4-week ahead')
ax[0].tick_params(labelrotation=45)
ax[0].legend()
ax[1].plot(actual_2024['dt'], actual_2024['Cases'], label = 'Actual Cases')
ax[1].plot(date, prediction_1_week, color = 'red', label = '1-week ahead')
ax[1].plot(date, prediction_4_week, color = 'green', label = '4-week ahead')
ax[1].tick_params(labelrotation=45)
ax[1].legend()
plt.suptitle('Comparing Actual and predicted')
plt.show()
```


Comparing Actual and predicted



[]: